

## Section III - SELECT

### 3.1: Selecting all columns

```
SELECT * FROM CUSTOMER;
```

To limit the number of records returned, use ROWNUM. To limit the results to just 2 records:

```
SELECT * FROM CUSTOMER WHERE ROWNUM <= 2;
```

### 3.2: Selecting specific columns

```
SELECT CUSTOMER_ID, NAME FROM CUSTOMER;
```

### 3.3: Expressions

First, select everything from PRODUCT

```
SELECT * FROM PRODUCT;
```

You can use expressions by declaring a TAXED\_PRICE . This is not a column, but rather something that is calculated every time this query is executed.

```
SELECT PRODUCT_ID,  
DESCRIPTION,  
PRICE,  
PRICE * 1.07 AS TAXED_PRICE  
FROM PRODUCT;
```

In Oracle SQL Developer, you can hit CTRL + SPACE to show an autocomplete box with available fields.

You can also use aliases to declare an UNTAXED\_PRICE column off the PRICE , without any expression.

```
SELECT PRODUCT_ID,  
DESCRIPTION,  
PRICE as UNTAXED_PRICE,
```

```
PRICE * 1.07 AS TAXED_PRICE  
FROM PRODUCT;
```

**SWITCH TO SLIDES FOR MATHEMATICAL OPERATORS**

## 3.4: Using `round()` Function

```
SELECT PRODUCT_ID,  
DESCRIPTION,  
PRICE,  
ROUND(PRICE * 1.07, 2) AS TAXED_PRICE  
  
FROM PRODUCT;
```

## 3.5: Text Concatenation

You can slap a dollar sign to our result using concatenation.

```
SELECT PRODUCT_ID,  
DESCRIPTION,  
PRICE AS UNTAXED_PRICE,  
'$' || ROUND(PRICE * 1.07, 2) AS TAXED_PRICE  
FROM PRODUCT
```

You can merge text via concatenation. For instance, you can concatenate two fields and put a comma and space , in between.

```
SELECT NAME,  
CITY || ', ' || STATE AS LOCATION  
FROM CUSTOMER;
```

You can concatenate several fields to create an address.

```
SELECT NAME,  
STREET_ADDRESS || ' ' || CITY || ', ' || STATE || ' ' || ZIP AS SHIP_ADDRESS  
FROM CUSTOMER;
```

This works with any data types, like numbers, texts, and dates. Also note that some platforms use `concat()` function instead of double pipes ||

**SWITCH TO SLIDES FOR EXERCISE**

## 3.6: Comments

To make a comments in SQL, use commenting dashes or blocks:

```
-- this is a comment  
  
/*  
This is a  
multiline comment  
*/
```

## Section IV- WHERE

### 4.1: Getting year 2010 records

```
SELECT * FROM station_data  
WHERE year = 2010;
```

### 4.2: Getting non-2010 records

```
SELECT * FROM station_data  
WHERE year != 2010;
```

```
SELECT * FROM station_data  
WHERE year <> 2010;
```

### 4.3: Getting records between 2005 and 2010

```
SELECT * FROM station_data  
WHERE year BETWEEN 2005 AND 2010
```

### 4.4: Using AND

```
SELECT * FROM station_data  
WHERE year >= 2005 AND year <= 2010
```

### 4.5: Exclusive Range

This will get the years between 2005 and 2010, but exclude 2005 and 2010

```
SELECT * FROM station_data  
WHERE year > 2005 AND year < 2010
```

## 4.6: Using OR

```
SELECT * FROM station_data  
WHERE MONTH = 3  
OR MONTH = 6  
OR MONTH = 9  
OR MONTH = 12
```

## 4.7: Using IN

```
SELECT * FROM station_data  
WHERE MONTH IN (3,6,9,12);
```

## 4.8: Using NOT IN

```
SELECT * FROM station_data  
WHERE MONTH NOT IN (3,6,9,12);
```

## 4.9: Using Modulus

The modulus will perform division but return the remainder. So a remainder of 0 means the two numbers divide evenly.

```
SELECT * FROM station_data  
WHERE MOD(MONTH, 3) = 0;
```

## 4.10: Using WHERE on TEXT

```
SELECT * FROM station_data  
WHERE report_code = '513A63'
```

## 4.11: Using IN with text

```
SELECT * FROM station_data  
WHERE report_code IN ('513A63', '1F8A7B', 'EF616A')
```

## 4.12: Using `length()` function

```
SELECT * FROM station_data  
WHERE LENGTH(report_code) != 6
```

## 4.13A: Using `LIKE` for any characters

```
SELECT * FROM station_data  
WHERE report_code LIKE 'A%';
```

## 4.13B: Using Regular Expressions

If you are familiar with regular expressions, you can use those to identify and qualify text patterns.

```
SELECT * FROM STATION_DATA  
WHERE REGEXP_LIKE(report_code, '^A.*$')
```

## 4.14: Using `LIKE` for one character

```
SELECT * FROM station_data  
WHERE report_code LIKE 'B_C%';
```

| For `LIKE` , `%` is used in a different context than modulus `%`

## 4.15: True Booleans 1

```
SELECT * FROM station_data  
WHERE tornado = 1 AND hail = 1;
```

## 4.16: True Booleans 2

```
SELECT * FROM station_data  
WHERE tornado = 1 AND hail = 1
```

## 4.17: False Booleans 1

```
SELECT * FROM station_data  
WHERE tornado = 0 AND hail = 1;
```

## 4.18: False Booleans 2

```
SELECT * FROM station_data  
WHERE tornado = 0 AND hail = 1;
```

## 4.19: Handling NULL

A NULL is an absent value. It is not zero, empty text '', or any value. It is blank.

To check for a null value:

```
SELECT * FROM station_data  
WHERE snow_depth IS NULL;
```

## 4.20: Handling NULL in conditions

Nulls will not qualify with any condition that doesn't explicitly handle it.

```
SELECT * FROM station_data  
WHERE precipitation <= 0.5;
```

If you want to include nulls, do this:

```
SELECT * FROM station_data  
WHERE precipitation IS NULL OR precipitation <= 0.5;
```

## 4.21: Combining AND and OR

Querying for sleet or snow

Problematic. What belongs to the AND and what belongs to the OR ?

```
SELECT * FROM station_data  
WHERE rain = 1 AND temperature <= 32  
OR snow_depth > 0;
```

You must group up the sleet condition in parenthesis so it is treated as one unit.

```
SELECT * FROM station_data  
WHERE (rain = 1 AND temperature <= 32)  
OR snow_depth > 0;
```

## Exercises

```
-- SELECT all records where TEMPERATURE is between 30 and 50 degrees
```

```
SELECT * FROM station_data  
WHERE temperature BETWEEN 30 AND 50;  
-- OR  
SELECT * FROM station_data  
WHERE temperature >= 30 and temperature <= 50;
```

```
-- SELECT all records where station_pressure is greater than 1000 and a  
tornado was present
```

```
SELECT * FROM STATION_DATA  
WHERE station_pressure > 1000 AND tornado = 1;  
-- OR  
SELECT * FROM STATION_DATA  
WHERE station_pressure > 1000 AND tornado = 1;
```

```
-- SELECT all records with report codes E6AED7, B950A1, 98DDAD
```

```
SELECT * FROM STATION_DATA  
WHERE report_code IN ('E6AED7', 'B950A1', '98DDAD')  
-- OR  
SELECT * FROM STATION_DATA  
WHERE report_code = 'E6AED7'  
OR report_code = 'B950A1'  
OR report_code = '98DDAD'
```

```
-- SELECT all records WHERE station_pressure is null
```

```
SELECT * FROM STATION_DATA  
WHERE station_pressure IS NULL;
```

## Section V- GROUP BY and ORDER BY

### 5.1: Getting a count of records

```
SELECT count(*) as record_count FROM station_data
```

### 5.2 Getting a count of records with a condition

```
SELECT count(*) as record_count FROM station_data
WHERE tornado = 1
```

## 5.3 Getting a count by year

```
SELECT year, count(*) as record_count
FROM station_data
WHERE tornado = 1
GROUP BY year
```

## 5.4 Getting a count by year, month

```
SELECT year, month, count(*) as record_count
FROM station_data
WHERE tornado = 1
GROUP BY year, month
```

## 5.5 Getting a count by year, month with ordinal index

```
-- not supported in ORACLE by default
SELECT year, month, count(*) as record_count
FROM station_data
WHERE tornado = 1
GROUP BY 1, 2
```

## 5.6 Using ORDER BY

```
SELECT year, month, count(*) as record_count
FROM station_data
WHERE tornado = 1
GROUP BY year, month
ORDER BY year, month
```

## 5.7 Using ORDER BY with DESC

```
SELECT year, month, count(*) as record_count
FROM station_data
WHERE tornado = 1
GROUP BY year, month
ORDER BY year DESC, month
```

## 5.8 Counting non-null values

```
SELECT COUNT(snow_depth) as recorded_snow_depth_count  
FROM station_data
```

## 5.9 Average temperature by month since year 2000

```
SELECT month, AVG(temperature) as avg_temp  
FROM station_data  
WHERE year >= 2000  
GROUP BY month
```

## 5.10 Average temperature (with rounding) by month since year 2000

```
SELECT month, ROUND(AVG(temperature),2) as avg_temp  
FROM station_data  
WHERE year >= 2000  
GROUP BY month
```

## 5.11 Sum of snow depth

```
SELECT year, SUM(snow_depth) as total_snow  
FROM station_data  
WHERE year >= 2005  
GROUP BY year
```

## 5.12 Multiple aggregations

```
SELECT year,  
SUM(snow_depth) as total_snow,  
SUM(precipitation) as total_precipitation,  
MAX(precipitation) as max_precipitation  
  
FROM station_data  
WHERE year >= 2005  
GROUP BY year
```

## 5.13 Using HAVING

You cannot use WHERE on aggregations. This will result in an error.

```
-- this will error, why?  
SELECT year,  
SUM(precipitation) as total_precipitation  
FROM station_data  
WHERE total_precipitation > 30  
GROUP BY year
```

You can however, use HAVING.

```
SELECT year,  
SUM(precipitation) as total_precipitation  
FROM station_data  
GROUP BY year  
HAVING total_precipitation > 30
```

Note that Oracle does not support aliasing in GROUP BY and HAVING.

Therefore you have to rewrite the entire expression each time

```
SELECT year,  
SUM(precipitation) as total_precipitation  
FROM station_data  
GROUP BY year  
HAVING SUM(precipitation) > 30
```

## 5.14 Getting Distinct values

You can get DISTINCT values for one or more columns

```
SELECT DISTINCT station_number FROM station_data
```

You can also get distinct combinations of values for multiple columns

```
SELECT DISTINCT station_number, year FROM station_data
```

## Exercise

```
-- Find the SUM of precipitation by year when a tornado was present, and  
sort by year descending.
```

```
SELECT year,  
SUM(precipitation) AS tornado_precipitation  
FROM station_data  
WHERE tornado = 1  
GROUP BY year  
ORDER BY year DESC
```

```
-- SELECT the year and max snow depth, but only years where the max snow  
depth is at least 50.
```

```
SELECT year,  
max(snow_depth) AS max_snow_depth  
FROM STATION_DATA  
GROUP BY year  
HAVING max(snow_depth) >= 50
```

## Section VI - CASE Statements

### 6.1 Categorizing Wind Speed

You can use a `CASE` statement to turn a column value into another value based on conditions. For instance, we can turn different `wind_speed` ranges into `HIGH`, `MODERATE`, and `LOW` categories.

```
SELECT report_code, year, month, day, wind_speed,  
  
CASE  
    WHEN wind_speed >= 40 THEN 'HIGH'  
    WHEN wind_speed >= 30 THEN 'MODERATE'  
    WHEN wind_speed >= 0 THEN 'LOW'  
    ELSE 'N/A'  
END AS wind_severity  
  
FROM station_data  
  
ORDER by wind_speed DESC
```

### 6.2 More Efficient Way To Categorize Wind Speed

We can actually omit `AND wind_speed < 40` from the previous example because each `WHEN / THEN` is evaluated from top-to-bottom. The first one it finds to be true is the one it will go with, and stop evaluating subsequent conditions.

```
SELECT report_code, year, month, day, wind_speed,  
  
CASE  
    WHEN wind_speed >= 40 THEN 'HIGH'  
    WHEN wind_speed >= 30 THEN 'MODERATE'  
    ELSE 'LOW'  
END AS wind_severity  
  
FROM station_data
```

## 6.3 Using CASE with GROUP BY

We can use `GROUP BY` in conjunction with a `CASE` statement to slice data in more ways, such as getting the record count by `wind_severity`.

```
SELECT  
  
CASE  
    WHEN wind_speed >= 40 THEN 'HIGH'  
    WHEN wind_speed >= 30 THEN 'MODERATE'  
    WHEN wind_speed >= 0 THEN 'LOW'  
    ELSE 'N/A'  
END AS wind_severity,  
  
COUNT(*) AS record_count  
  
FROM station_data  
  
GROUP BY wind_severity
```

Also, some `wind_speed` values are `NULL`, so without an `ELSE` any records that do not meet a condition will turn out to be `NULL`.

```
SELECT  
  
CASE  
    WHEN wind_speed >= 40 THEN 'HIGH'  
    WHEN wind_speed >= 30 THEN 'MODERATE'  
    WHEN wind_speed >= 0 THEN 'LOW'  
END AS wind_severity,  
  
COUNT(*) AS record_count  
  
FROM station_data
```

```
GROUP BY wind_severity
```

## 6.4 "Zero/Null" Case Trick

There is really no way to create multiple aggregations with different conditions unless you know a trick with the `CASE` statement. If you want to find two total precipitation, with and without tornado precipitations, for each year and month, you have to do separate queries.

### Tornado Precipitation

```
SELECT year, month,
SUM(precipitation) as tornado_precipitation
FROM station_data
WHERE tornado = 1
AND year >= 1990
GROUP BY year, month
```

### Non-Tornado Precipitation

```
SELECT year, month,
SUM(precipitation) as non_tornado_precipitation
FROM station_data
WHERE tornado = 0
AND year >= 1990
GROUP BY year, month
```

But you can use a single query using a `CASE` statement that sets a value to 0 if the condition is not met. That way it will not impact the sum.

```
SELECT year, month,
SUM(CASE WHEN tornado = 1 THEN precipitation ELSE 0 END) as
tornado_precipitation,
SUM(CASE WHEN tornado = 0 THEN precipitation ELSE 0 END) as
non_tornado_precipitation

FROM station_data
WHERE year >= 1990

GROUP BY year, month
```

Many folks who are not aware of the zero/null case trick will resort to derived tables (not covered in this class but covered in *Advanced SQL for Data Analysis*), which adds an

unnecessary amount of effort and mess.

```
SELECT t.year,
t.month,
t.tornado_precipitation,
non_t.non_tornado_precipitation

FROM (
    SELECT year, month,
    SUM(precipitation) as tornado_precipitation
    FROM station_data
    WHERE tornado = 1
    AND year >= 1990
    GROUP BY year, month
) t

INNER JOIN

(
    SELECT year, month,
    SUM(precipitation) as non_tornado_precipitation
    FROM station_data
    WHERE tornado = 0
    AND year >= 1990
    GROUP BY year, month
) non_t
ON t.year = non_t.year AND t.month = non_t.month
```

## 6.5 Using Null in a CASE to conditionalize MIN/MAX

Since NULL is ignored in SUM, MIN, MAX, and other aggregate functions, you can use it in a CASE statement to conditionally control whether or not a value should be included in that aggregation.

For instance, we can split up max precipitation when a tornado was present vs not present.

```
SELECT year,
MAX(CASE WHEN tornado = 0 THEN precipitation ELSE NULL END) as
max_non_tornado_precipitation,
MAX(CASE WHEN tornado = 1 THEN precipitation ELSE NULL END) as
max_tornado_precipitation
FROM station_data
WHERE year >= 1990
GROUP BY year
```

*Switch to slides for exercise*

## Exercise 6.1

SELECT the report\_code, year, quarter, and temperature, where a “quarter” is “Q1”, “Q2”, “Q3”, or “Q4” reflecting months 1-3, 4-6, 7-9, and 10-12 respectively.

**ANSWER:**

```
SELECT
    report_code,
    year,
    CASE
        WHEN month BETWEEN 1 and 3 THEN 'Q1'
        WHEN month BETWEEN 4 and 6 THEN 'Q2'
        WHEN month BETWEEN 7 and 9 THEN 'Q3'
        WHEN month BETWEEN 10 and 12 THEN 'Q4'
    END as quarter,
    temperature
FROM STATION_DATA
```

## Exercise 6.2

Get the average temperature by quarter and year, where a “quarter” is “Q1”, “Q2”, “Q3”, or “Q4” reflecting months 1-3, 4-6, 7-9, and 10-12 respectively.

**ANSWER**

```
SELECT
    year,
    CASE
        WHEN month BETWEEN 1 and 3 THEN 'Q1'
        WHEN month BETWEEN 4 and 6 THEN 'Q2'
        WHEN month BETWEEN 7 and 9 THEN 'Q3'
        WHEN month BETWEEN 10 and 12 THEN 'Q4'
    END as quarter,
    AVG(temperature) as avg_temp
```

```
FROM STATION_DATA  
GROUP BY year, quarter
```

## Section VII - JOIN

### 7.1A INNER JOIN

(Refer to slides Section VII)

View customer address information with each order by joining tables CUSTOMER and CUSTOMER\_ORDER .

```
SELECT ORDER_ID,  
CUSTOMER.CUSTOMER_ID,  
ORDER_DATE,  
SHIP_DATE,  
NAME,  
STREET_ADDRESS,  
CITY,  
STATE,  
ZIP,  
PRODUCT_ID,  
ORDER_QTY  
  
FROM CUSTOMER INNER JOIN CUSTOMER_ORDER  
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID
```

Joins allow us to keep stored data normalized and simple, but we can get more descriptive views of our data by using joins.

Notice how two customers are omitted since they don't have any orders (refer to slides).

### 7.2B A BAD APPROACH

You may come across a style of joining where commas are used to select the needed tables, and a WHERE defines the join condition as shown below:

```
SELECT ORDER_ID,  
CUSTOMER.CUSTOMER_ID,  
ORDER_DATE,  
SHIP_DATE,  
NAME,  
STREET_ADDRESS,  
CITY,
```

```
STATE,  
ZIP,  
PRODUCT_ID,  
ORDER_QTY  
  
FROM CUSTOMER, CUSTOMER_ORDER  
WHERE CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID
```

Do not use this approach no matter how much your colleagues use it (and educate them not to use it either). It is extremely inefficient as it will generate a cartesian product across both tables (every possible combination of records between both), and then filter it based on the WHERE. It does not work with `LEFT JOIN` either, which we will look at shortly.

Using the `INNER JOIN` with an `ON` condition avoids the cartesian product and is more efficient. Therefore, always use that approach.

## 7.2 LEFT OUTER JOIN

To include all customers, regardless of whether they have orders, you can use a left outer join via `LEFT JOIN` (refer to slides).

If any customers do not have any orders, they will get one record where the `CUSTOMER_ORDER` fields will be null.

```
SELECT CUSTOMER.CUSTOMER_ID,  
NAME,  
STREET_ADDRESS,  
CITY,  
STATE,  
ZIP,  
ORDER_DATE,  
SHIP_DATE,  
ORDER_ID,  
PRODUCT_ID,  
ORDER_QTY  
  
FROM CUSTOMER LEFT JOIN CUSTOMER_ORDER  
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID
```

## 7.3 Finding Customers with No Orders

With a left outer join, you can filter for NULL values on the `CUSTOMER_ORDER` table to find customers that have no orders.

```
SELECT CUSTOMER.CUSTOMER_ID,
NAME AS CUSTOMER_NAME

FROM CUSTOMER LEFT JOIN CUSTOMER_ORDER
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID

WHERE ORDER_ID IS NULL
```

You can use a left outer join to find child records with no parent, or parent records with no children (e.g. a CUSTOMER\_ORDER with no CUSTOMER , or a CUSTOMER with no CUSTOMER\_ORDER s).

## 7.4 Joining Multiple Tables

Bring in PRODUCT to supply product information for each CUSTOMER\_ORDER , on top of CUSTOMER information.

```
SELECT ORDER_ID,
CUSTOMER.CUSTOMER_ID,
NAME AS CUSTOMER_NAME,
STREET_ADDRESS,
CITY,
STATE,
ZIP,
ORDER_DATE,
PRODUCT.PRODUCT_ID,
DESCRIPTION,
ORDER_QTY

FROM CUSTOMER INNER JOIN CUSTOMER_ORDER
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID

INNER JOIN PRODUCT
ON CUSTOMER_ORDER.PRODUCT_ID = PRODUCT.PRODUCT_ID
```

## 7.7 Using Expressions with JOINs

You can use expressions combining any fields on any of the joined tables. For instance, we can now get the total revenue for each customer.

```
SELECT ORDER_ID,
CUSTOMER.CUSTOMER_ID,
NAME AS CUSTOMER_NAME,
```

```
STREET_ADDRESS,  
CITY,  
STATE,  
ZIP,  
ORDER_DATE,  
PRODUCT.PRODUCT_ID,  
DESCRIPTION,  
ORDER_QTY,  
ORDER_QTY * PRICE AS REVENUE  
  
FROM CUSTOMER INNER JOIN CUSTOMER_ORDER  
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID  
  
INNER JOIN PRODUCT  
ON CUSTOMER_ORDER.PRODUCT_ID = PRODUCT.PRODUCT_ID
```

## 7.6 Using GROUP BY with JOINS

You can use `GROUP BY` with a join. For instance, you can find the total revenue for each customer by leveraging all three joined tables, and aggregating the `REVENUE` expression we created earlier.

```
SELECT  
CUSTOMER.CUSTOMER_ID,  
NAME AS CUSTOMER_NAME,  
sum(ORDER_QTY * PRICE) AS TOTAL_REVENUE  
  
FROM CUSTOMER INNER JOIN CUSTOMER_ORDER  
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID  
  
INNER JOIN PRODUCT  
ON CUSTOMER_ORDER.PRODUCT_ID = PRODUCT.PRODUCT_ID  
  
GROUP BY CUSTOMER.CUSTOMER_ID, NAME
```

To see all customers even if they had no orders, use a `LEFT JOIN`

```
SELECT  
CUSTOMER.CUSTOMER_ID,  
NAME AS CUSTOMER_NAME,  
sum(ORDER_QTY * PRICE) AS TOTAL_REVENUE  
  
FROM CUSTOMER LEFT JOIN CUSTOMER_ORDER  
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID
```

```
LEFT JOIN PRODUCT
ON CUSTOMER_ORDER.PRODUCT_ID = PRODUCT.PRODUCT_ID

GROUP BY CUSTOMER.CUSTOMER_ID, NAME
```

You can also use a `NVL()` function to turn null sums into zeros.

```
SELECT
CUSTOMER.CUSTOMER_ID,
NAME AS CUSTOMER_NAME,
NVL(sum(ORDER_QTY * PRICE), 0) AS TOTAL_REVENUE

FROM CUSTOMER LEFT JOIN CUSTOMER_ORDER
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID

LEFT JOIN PRODUCT
ON CUSTOMER_ORDER.PRODUCT_ID = PRODUCT.PRODUCT_ID

GROUP BY CUSTOMER.CUSTOMER_ID, NAME
```

## Exercise

```
/*
SELECT the ORDER_ID, ORDER_DATE, and DESCRIPTION (from PRODUCT)
(hint, you will need to INNER JOIN CUSTOMER_ORDER and PRODUCT)
*/
SELECT ORDER_ID, ORDER_DATE, DESCRIPTION

FROM CUSTOMER_ORDER INNER JOIN PRODUCT
ON CUSTOMER_ORDER.PRODUCT_ID = PRODUCT.PRODUCT_ID
```

```
/*
Find the total revenue by product. Include the fields PRODUCT_ID,
DESCRIPTION, and then the TOTAL_REVENUE.

(Hint: you will need to join CUSTOMER_ORDER and PRODUCT. Then do a GROUP BY)
*/
SELECT PRODUCT.PRODUCT_ID,
DESCRIPTION,
NVL(SUM (ORDER_QTY * PRICE), 0) AS TOTAL_REVENUE

FROM PRODUCT LEFT JOIN CUSTOMER_ORDER
```

```
ON PRODUCT.PRODUCT_ID = CUSTOMER_ORDER.PRODUCT_ID  
GROUP BY PRODUCT.PRODUCT_ID, DESCRIPTION
```

## Section VIII - Database Design

Refer to slides for database design concepts

To view source code for SQL Injection Demo, here is the GitHub page:

<https://github.com/thomasnield/sql-injection-demo>

To read about normalized forms (which we do not cover in favor of a more intuitive approach), you can read this article:

<http://www.dummies.com/programming/sql/sql-first-second-and-third-normal-forms/>

### 7.1 - Creating a Table

In Oracle, you can use SQL Developer or command line tools to create a new schema or user, and then create tables.

Create the ATTENDEE table. You can execute the following SQL.

```
CREATE TABLE ATTENDEE (  
    ATTENDEE_ID NUMBER PRIMARY KEY,  
    FIRST_NAME VARCHAR2 (30) NOT NULL,  
    LAST_NAME VARCHAR2 (30) NOT NULL,  
    PHONE NUMBER,  
    EMAIL VARCHAR2 (30),  
    VIP NUMBER(1) DEFAULT 0  
) ;
```

After each field declaration, we create "rules" for that field. For example, COMPANY\_ID must be a NUMBER , it is a PRIMARY KEY . The NAME field holds text because it is VARCHAR2 (a variable number of characters), and it is limited to 30 characters and cannot be NULL .

Lastly, we declare any FOREIGN KEY constraints, specifying which field is a FOREIGN KEY and what PRIMARY KEY it references. In this example, PRIMARY\_CONTACT\_ATTENDEE\_ID "references" the ATTENDEE\_ID in the ATTENDEE table, and it can only be those values.

### 7.2 - Creating the other tables

Create the other tables by executing the following SQL code.

```

CREATE TABLE COMPANY (
    COMPANY_ID NUMBER PRIMARY KEY,
    NAME VARCHAR2(30) NOT NULL,
    DESCRIPTION VARCHAR2(60),
    PRIMARY_CONTACT_ATTENDEE_ID NUMBER NOT NULL,
    CONSTRAINT FK_COMPANY_ATTENDEE FOREIGN KEY (PRIMARY_CONTACT_ATTENDEE_ID)
    REFERENCES ATTENDEE(ATTENDEE_ID)
);

CREATE TABLE ROOM (
    ROOM_ID NUMBER PRIMARY KEY,
    FLOOR_NUMBER NUMBER NOT NULL,
    SEAT_CAPACITY NUMBER NOT NULL
);

CREATE TABLE PRESENTATION (
    PRESENTATION_ID NUMBER PRIMARY KEY,
    BOOKED_COMPANY_ID NUMBER NOT NULL,
    BOOKED_ROOM_ID NUMBER NOT NULL,
    START_TIME TIMESTAMP,
    END_TIME TIMESTAMP,
    CONSTRAINT FK_PRESENTATION_COMPANY FOREIGN KEY (BOOKED_COMPANY_ID)
    REFERENCES COMPANY(COMPANY_ID),
    CONSTRAINT FK_PRESENTATION_ROOM FOREIGN KEY (BOOKED_ROOM_ID) REFERENCES
    ROOM(ROOM_ID)
);

CREATE TABLE PRESENTATION_ATTENDANCE (
    TICKET_ID NUMBER PRIMARY KEY,
    PRESENTATION_ID NUMBER,
    ATTENDEE_ID NUMBER,
    CONSTRAINT FK_PA_PRESENTATION FOREIGN KEY (PRESENTATION_ID) REFERENCES
    PRESENTATION(PRESENTATION_ID),
    CONSTRAINT FK_PA_ATTENDEE FOREIGN KEY (ATTENDEE_ID) REFERENCES
    ATTENDEE(ATTENDEE_ID)
);

```

## Creating Views

It is not uncommon to save `SELECT` queries that are used frequently into a database. These are known as **Views** and act very similarly to tables. You can essentially save a `SELECT` query and work with it just like a table.

For instance, say we wanted to save this SQL query that includes `ROOM` and `COMPANY` info with each `PRESENTATION` record.

```
SELECT COMPANY.NAME as BOOKED_COMPANY,
ROOM.ROOM_ID as ROOM_NUMBER,
ROOM.FLOOR_NUMBER as FLOOR,
ROOM.SEAT_CAPACITY as SEATS,
START_TIME, END_TIME

FROM PRESENTATION

INNER JOIN COMPANY
ON PRESENTATION.BOOKED_COMPANY_ID = COMPANY.COMPANY_ID

INNER JOIN ROOM
ON PRESENTATION.BOOKED_ROOM_ID = ROOM.ROOM_ID
```

You can save this as a view by executing the following SQL syntax: `CREATE VIEW [view name] AS [a SELECT query]`. For this example, this is what it would look like.

```
CREATE VIEW PRESENTATION_VW AS

SELECT COMPANY.NAME as BOOKED_COMPANY,
ROOM.ROOM_ID as ROOM_NUMBER,
ROOM.FLOOR_NUMBER as FLOOR,
ROOM.SEAT_CAPACITY as SEATS,
START_TIME, END_TIME

FROM PRESENTATION

INNER JOIN COMPANY
ON PRESENTATION.BOOKED_COMPANY_ID = COMPANY.COMPANY_ID

INNER JOIN ROOM
ON PRESENTATION.BOOKED_ROOM_ID = ROOM.ROOM_ID
```

You can then query it just like a table.

```
SELECT * FROM PRESENTATION_VW
WHERE SEATS >= 30
```

Obviously, there is no data yet so you will not get any results. But there will be once you populate data into this database.

## Section IX - Writing Data

In this section, we will learn how to write, modify, and delete data in a database.

## 9.1 Using INSERT

To create a new record in a table, use the `INSERT` command and supply the values for the needed columns.

Put yourself into the `ATTENDEE` table.

```
INSERT INTO ATTENDEE (ATTENDEE_ID, FIRST_NAME, LAST_NAME)
VALUES (1, 'Thomas', 'Nield')
```

Notice above that we declare the table we are writing to, which is `ATTENDEE`. We indicate the fields followed by the values. The `PHONE`, `EMAIL`, and `VIP` fields have default values or are nullable, and therefore optional.

## 9.2 Multiple INSERT records

You can insert multiple rows in an `INSERT`. This will add three people to the `ATTENDEE` table.

```
INSERT INTO ATTENDEE (ATTENDEE_ID, FIRST_NAME, LAST_NAME, PHONE, EMAIL, VIP)
VALUES (2, 'Jon', 'Skeeter', 4802185842, 'john.skeeter@rex.net', 1),
       (3, 'Sam', 'Scala', 2156783401, 'sam.scala@gmail.com', 0),
       (4, 'Brittany', 'Fisher', 5932857296, 'brittany.fisher@outlook.com', 0)
```

## 9.3 Testing the foreign keys

Let's test our design and make sure our primary/foreign keys are working.

Try to `INSERT` a `COMPANY` with a `PRIMARY_CONTACT_ATTENDEE_ID` that does not exist in the `ATTENDEE` table.

```
INSERT INTO COMPANY (COMPANY_ID, NAME, DESCRIPTION,
PRIMARY_CONTACT_ATTENDEE_ID)
VALUES (1, 'RexApp Solutions', 'A mobile app delivery service', 5)
```

Currently, there is no `ATTENDEE` with an `ATTENDEE_ID` of 5, this should error out which is good. It means we kept bad data out.

If you use an `ATTENDEE_ID` value that does exist and supply it as a `PRIMARY_CONTACT_ATTENDEE_ID`, we should be good to go.

```
INSERT INTO COMPANY (COMPANY_ID, NAME, DESCRIPTION,  
PRIMARY_CONTACT_ATTENDEE_ID)  
VALUES (1, 'RexApp Solutions', 'A mobile app delivery service', 3)
```

## 9.3 DELETE records

The `DELETE` command is dangerously simple. To delete records from both the `COMPANY` and `ATTENDEE` tables, execute the following SQL commands.

```
DELETE FROM COMPANY;  
DELETE FROM ATTENDEE;
```

Note that the `COMPANY` table has a foreign key relationship with the `ATTENDEE` table. Therefore we will have to delete records from `COMPANY` first before it allows us to delete data from `ATTENDEE`. Otherwise we will get a "FOREIGN KEY constraint failed effort" due to the `COMPANY` record we just added which is tied to the `ATTENDEE` with the `ATTENDEE_ID` of 3.

You can also use a `WHERE` to only delete records that meet a conditional. To delete all `ATTENDEE` records with no `PHONE` or `EMAIL`, you can run this command.

```
DELETE FROM ATTENDEE  
WHERE PHONE IS NULL AND EMAIL IS NULL
```

A good practice is to use a `SELECT *` in place of the `DELETE` first. That way you can get a preview of what records will be deleted with that `WHERE` condition.

```
SELECT * FROM ATTENDEE  
WHERE PHONE IS NULL AND EMAIL IS NULL
```

## UPDATE records

Say we wanted to change the phone number for the `ATTENDEE` with the `ATTENDEE_ID` value of 3, which is Sam Scala. We can do this with an `UPDATE` statement.

```
UPDATE ATTENDEE SET PHONE = 4802735872  
WHERE ATTENDEE_ID = 3
```

Using a `WHERE` is important, otherwise it will update all records with the specified `SET` assignment. This can be handy if you wanted to say, make all `EMAIL` values uppercase.

```
UPDATE ATTENDEE SET EMAIL = UPPER(EMAIL)
```

## 9.4 Dropping Tables

If you want to delete a table, it also is dangerously simple. Be very careful and sure before you delete any table, because it will remove it permanently.

```
DROP TABLE MY_UNWANTED_TABLE
```

## 9.5 Transactions

Transactions are helpful when you want a series of writes to succeed.

Below, we execute two successful write operations within a transaction.

```
BEGIN  
  
INSERT INTO ROOM (ROOM_ID, FLOOR_NUMBER, SEAT_CAPACITY) VALUES (1, 9, 80);  
  
INSERT INTO ROOM (ROOM_ID, FLOOR_NUMBER, SEAT_CAPACITY) VALUES (2, 10, 110);  
  
-- Commit the transaction to make changes permanent  
  
COMMIT;  
  
END;
```

But if we ever encountered a failure with our write operations, we can call `ROLLBACK` instead of `COMMIT` to go back to the database state when `BEGIN` was called.

Below, we have a failed operation due to a broken `INSERT`.

```
BEGIN  
  
INSERT INTO ROOM (ROOM_ID, FLOOR_NUMBER, SEAT_CAPACITY) VALUES (3, 9, 80);  
  
INSERT INTO ROOM (ROOM_ID, FLOOR_NUMBER, SEAT_CAPACITY) VALUES (10, 110);  
  
-- Commit the transaction to make changes permanent  
  
COMMIT;
```

```
END;
```

This will automatically roll back, but we can also do a manual ROLLBACK to the last transaction at any time BEGIN was called with this.

```
ROLLBACK;
```

## 9.6 Creating Indexes

You can create an index on a certain column to speed up SELECT performance, such as the EMAIL column on the ATTENDEE table.

```
CREATE INDEX email_index ON ATTENDEE(EMAIL);
```

You can also create an index for a column that has unique values, and it will make a special optimization for that case.

```
CREATE UNIQUE INDEX email_index ON ATTENDEE(EMAIL);
```

To remove an index, use the DROP command.

```
DROP INDEX email_index;
```

## 9.7 Working with Dates and Times

Use the ISO 'yyyy-mm-dd' syntax with TO\_DATE to treat them as dates easily.

Keep in mind much of this functionality is proprietary to Oracle. Make sure you learn the date and time functionality for your specific database platform.

```
SELECT * FROM CUSTOMER_ORDER
WHERE SHIP_DATE < TO_DATE('2015-05-21', 'YYYY-MM-DD')
```

To get today's date, we use the SYSDATE function. The DUAL keyword is a single row/column dummy table when we don't want to retrieve data from any tables for demonstration in situations like this.

```
SELECT SYSDATE FROM DUAL
```

To shift a date:

```
SELECT SYSDATE - 1 FROM DUAL;
SELECT TO_DATE('2015-12-07', 'YYYY-MM-DD') + INTERVAL '3' MONTH - INTERVAL
'1' DAY FROM DUAL;
```

To work with times, use `hh:mm:ss` format.

```
SELECT TO_TIMESTAMP('16:31', 'HH24:MI') < TO_TIMESTAMP('08:31', 'HH24:MI')
FROM DUAL
```

To get today's GMT time:

```
SELECT CURRENT_TIMESTAMP FROM DUAL
```

To shift a time:

```
SELECT TO_TIMESTAMP('16:31', 'HH24:MI') + INTERVAL '1' MINUTE FROM DUAL
```

To merge a date and time, use a `TIMESTAMP` type.

```
SELECT TO_TIMESTAMP('2015-12-13 16:04:11', 'YYYY-MM-DD HH24:MI:SS') FROM
DUAL
SELECT TO_TIMESTAMP('2015-12-13 16:04:11', 'YYYY-MM-DD HH24:MI:SS') -
INTERVAL '1' DAY + INTERVAL '3' HOUR FROM DUAL
```

To format dates and times a certain way:

```
SELECT TO_CHAR(SYSDATE, 'DD-MM-YYYY') FROM DUAL
```

Refer to Oracle documentation

[https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/TO\\_CHAR-datetime.html](https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/TO_CHAR-datetime.html)

Another helpful tutorial on using dates and times with Oracle.

<https://www.oracletutorial.com/oracle-basics/oracle-date/>